

# Coursera Peer-graded Assignment: Deep Learning and Reinforcement Learning

## Objective: Object Classification and Localization

For an object classification and localization problem, we start off using the same network we saw in image classification.

So, we have an image as an input, which goes through a ConvNet that results in a vector of features fed to a softmax to classify the object (for example with 4 classes for pedestrians/cars/bike/background).

Now, if we want to localize those objects in the image as well, we change the neural network to have a few more output units that encompass a bounding box.

In particular, we add four more numbers, which identify the x and y coordinates of the lower left corner and upper right corner of the box ( $x_{min}$ ,  $y_{min}$ ,  $x_{max}$ ,  $y_{max}$ ).

## Dataset: The Oxford-IIIT Pet Dataset

A 37 category pet dataset with roughly 200 images for each class.

The images have a large variations in scale, pose and lighting.

All images have an associated ground truth annotation of breed, head ROI, and pixel level trimap segmentation.

<https://www.robots.ox.ac.uk/~vgg/data/pets/>

Omkar M Parkhi and Andrea Vedaldi and Andrew Zisserman and C. V. Jawahar

In [1]:

```
#import wget
```

```
#url='https://www.robots.ox.ac.uk/~vgg/data/pets/data/images.tar.gz'  
#wget.download(url,'images.tar.gz')  
#url='https://www.robots.ox.ac.uk/~vgg/data/pets/data/annotations.tar.gz'  
#wget.download(url,'annotations.tar.gz')
```

In [2]:

```
#images_tar=tarfile.open('images.tar.gz')  
#images_tar.extractall('./Oxford_Data')  
#images_tar.close()
```

In [3]:

```
#annotations_tar=tarfile.open('annotations.tar.gz')  
#annotations_tar.extractall('./Oxford_Data')  
#annotations_tar.close()
```

## Importing necessary libraries

In [4]:

```
import os  
from PIL import Image,ImageOps  
import numpy as np  
from collections import namedtuple  
import matplotlib.pyplot as plt  
from matplotlib.patches import Rectangle  
import xml.etree.ElementTree as ET  
import pandas as pd  
import tensorflow as tf  
from tensorflow.keras.applications.xception import Xception,preprocess_input  
from tensorflow.keras.models import Model,Sequential  
from tensorflow.keras.layers import Dense, GlobalAveragePooling2D, Dropout, Flatten  
import tarfile
```

## Directory Structure

In [5]:

```
os.getcwd()
```

Out[5]:

```
'C:\\\\Users\\\\Asus\\\\Downloads'
```

```
In [6]: os.listdir(os.path.join(os.getcwd(), 'Oxford_Data'))
```

```
Out[6]: ['annotations', 'images']
```

```
In [7]: len(os.listdir(os.path.join(os.getcwd(), 'Oxford_Data\\images')))
```

```
Out[7]: 7392
```

```
In [8]: os.listdir(os.path.join(os.getcwd(), 'Oxford_Data\\annotations'))
```

```
Out[8]: ['.trimaps',
          'list.txt',
          'README',
          'test.txt',
          'trainval.txt',
          'trimaps',
          'xmls']
```

xmls folder contains bounding box attributes corresponding to trainval data

```
In [9]: all_images_path=os.path.join(os.getcwd(), 'Oxford_Data\\images')
all_xmls_path=os.path.join(os.getcwd(), 'Oxford_Data\\annotations\\xmls')
```

## 'process\_image(args...)': to rescale the image and corresponding bounding box

```
In [10]: Bounding_Box=namedtuple('BoundingBox', ['xmin', 'ymin', 'xmax', 'ymax'])
```

```
def process_image(path_to_image, bounding_box=None):
```

```
    target_size=(224,224)
```

```
    img=Image.open(path_to_image)
```

```
    # Padding image to make it Square
```

```
    width,height=img.size
```

```
    h_pad=0
```

```
    w_pad=0
```

```
    bonus_h_pad=0
```

```
bonus_w_pad=0
pix_diff=abs(width-height)
if width>height:
    h_pad=pix_diff//2
    bonus_h_pad=pix_diff%2
elif height>width:
    w_pad=pix_diff//2
    bonus_w_pad=pix_diff%2
img=ImageOps.expand(img,(w_pad,h_pad,w_pad+bonus_w_pad,h_pad+bonus_h_pad))

# Resizing the image

width,height=img.size
img=img.resize((target_size[0],target_size[1]))
img_data=np.array(img.getdata()).reshape(target_size[0],target_size[1],3)

# Rescaling Bounding Box Dimensions

if bounding_box is not None:

    width_scaling_factor=target_size[0]/width
    height_scaling_factor=target_size[1]/height

    new_xmin=bounding_box.xmin+w_pad
    new_xmin*=width_scaling_factor

    new_xmax=bounding_box.xmax+w_pad
    new_xmax*=width_scaling_factor

    new_ymin=bounding_box.ymin+h_pad
    new_ymin*=height_scaling_factor

    new_ymax=bounding_box.ymax+h_pad
    new_ymax*=height_scaling_factor

    return (img_data,Bounding_Box(new_xmin,new_ymin,new_xmax,new_ymax))

return (img_data,None)
```

'get\_bounding\_box(args...)':to get bounding box dimensions from xml file

In [11]:

```
def get_bounding_box(path_to_xml):

    tree=ET.parse(path_to_xml)
    root=tree.getroot()
    xmin=int(root.find('./object/bndbox/xmin').text)
    ymin=int(root.find('./object/bndbox/ymin').text)
    xmax=int(root.find('./object/bndbox/xmax').text)
    ymax=int(root.find('./object/bndbox/ymax').text)
    bnd_box=Bounding_Box(xmin,ymin,xmax,ymax)

    return bnd_box
```

## 'image\_with\_bndbox(args...)': to show an image with true and/or predicted bounded box

In [12]:

```
def image_with_bndbox(image_data,true_bndbox=None,pred_bndbox=None):

    fig,ax=plt.subplots(1)
    ax.imshow(image_data)

    # Drawing True Bounding Box over image

    if true_bndbox is not None:

        anchor_x=true_bndbox.xmin
        anchor_y=true_bndbox.ymin
        r1=Rectangle((anchor_x,anchor_y),width=true_bndbox.xmax-true_bndbox.xmin,height=true_bndbox.ymax-true_bndbox.ymin,color='g'
        ax.add_patch(r1)

    # Drawing Predicted Bounding Box over image

    if pred_bndbox is not None:

        r2=Rectangle((pred_bndbox.xmin,pred_bndbox.ymin),width=pred_bndbox.xmax-pred_bndbox.xmin,height=pred_bndbox.ymax-pred_bndb
        ax.add_patch(r2)

    # Calculating IOU

    if true_bndbox is not None and pred_bndbox is not None:

        xmin_inter=max(true_bndbox.xmin,pred_bndbox.xmin)
```

```
ymin_inter=max(true_bndbox.ymin,pred_bndbox.ymin)
xmax_inter=min(true_bndbox.xmax,pred_bndbox.xmax)
ymax_inter=min(true_bndbox.ymax,pred_bndbox.ymax)
true_bndbox_area=(true_bndbox.xmax-true_bndbox.xmin+1)*(true_bndbox.ymax-true_bndbox.ymin+1)
pred_bndbox_area=(pred_bndbox.xmax-pred_bndbox.xmin+1)*(pred_bndbox.ymax-pred_bndbox.ymin+1)
inter_area=max(0,xmax_inter-xmin_inter+1)*max(0,ymax_inter-ymax_inter+1)
iou=(inter_area)/float(true_bndbox_area+pred_bndbox_area-inter_area)
print('IOU:',iou)

plt.show()
```

## 'final\_image(args...)':acts as pipeline function to process and display image with boundng box

In [13]:

```
def final_image(image_name,pred_bndbox=None):

    image_path=os.path.join(all_images_path,image_name+'.jpg')
    xml_path=os.path.join(all_xmls_path,image_name+'.xml')
    true_bndbox=get_bounding_box(xml_path)
    img,true_bndbox=process_image(image_path,true_bndbox)
    image_with_bndbox(img,true_bndbox,pred_bndbox)
```

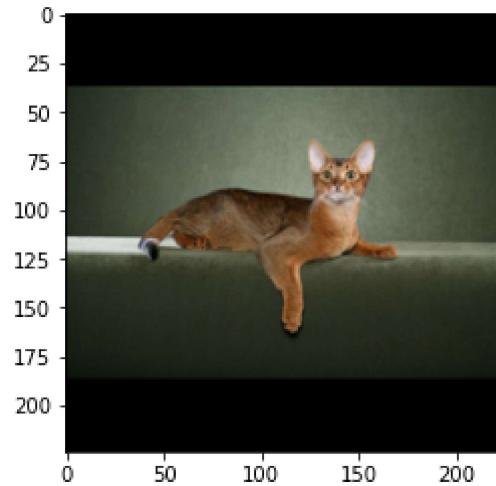
In [14]:

```
img,_=process_image('C:\\\\Users\\\\Asus\\\\Downloads\\\\Oxford_Data\\\\images\\\\Abyssinian_1.jpg')
```

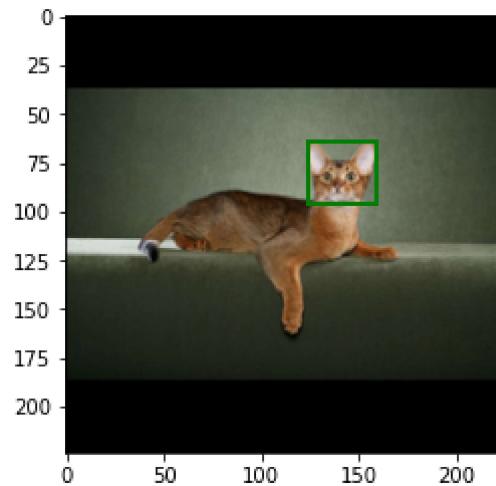
In [15]:

```
plt.imshow(img)
```

Out[15]:



```
In [16]: final_image('Abyssinian_1')
```



## Processing Training and Testing Data

### Training

```
In [17]: df_train=pd.read_csv(os.path.join(os.getcwd(),'Oxford_Data\\annotations\\trainval.txt'),sep=' ',names=['Image','ID','Species','Bre
```

```
In [18]: df_train.head()
```

Out[18]:

	Image	ID	Species	Breed
0	Abyssinian_100	1	1	1
1	Abyssinian_101	1	1	1
2	Abyssinian_102	1	1	1
3	Abyssinian_103	1	1	1
4	Abyssinian_104	1	1	1

```
In [19]: df_train['Species']=df_train['Species']-1
```

```
In [20]: print('No. of Cats:',df_train['Species'].value_counts()[0])
print('No. of Dogs:',df_train['Species'].value_counts()[1])
```

No. of Cats: 1188

No. of Dogs: 2492

```
In [21]: print('No. of Cat Species:',len(df_train.groupby('Species')['Breed'].unique()[0]))
print('No. of Dog Species:',len(df_train.groupby('Species')['Breed'].unique()[1]))
```

No. of Cat Species: 12

No. of Dog Species: 25

```
In [22]: train_tuple_df=np.array(df_train[['Image','Species']])
train_tuple_df
```

Out[22]:

```
array([['Abyssinian_100', 0],
       ['Abyssinian_101', 0],
       ['Abyssinian_102', 0],
       ...,
       ['yorkshire_terrier_189', 1],
       ['yorkshire_terrier_18', 1],
       ['yorkshire_terrier_190', 1]], dtype=object)
```

## 'process\_train\_data(args...)':to process and prepare training data

In [23]:

```
def process_train_data(df):

    data=list()

    for ele in df:

        image_name=ele[0]
        image_class=ele[1]

        if image_name+'.xml' not in os.listdir(all_xmls_path):

            continue

        else:

            image_path=os.path.join(all_images_path,image_name+'.jpg')
            xml_path=os.path.join(all_xmls_path,image_name+'.xml')
            true_bndbox=get_bounding_box(xml_path)
            img,true_bndbox=process_image(image_path,true_bndbox)
            data.append((image_name,img,true_bndbox,image_class))

    print('No. of Elements processed:',len(data))

    return np.array(data)
```

In [24]:

```
train_data=process_train_data(train_tuple_df)
```

No. of Elements processed: 3671

C:\Users\Asus\AppData\Local\Temp\ipykernel\_37160\3957214975.py:24: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

```
return np.array(data)
```

**train\_data[0]**-->image\_name

**train\_data[1]**-->img

**train\_data[2]**-->true\_bndbox

```
train_data[3]-->image_class
```

## Testing

```
In [25]: df_test=pd.read_csv(os.path.join(os.getcwd(),'Oxford_Data\\annotations\\test.txt'),sep=' ',names=['Image','ID','Species','Breed'])
```

```
In [26]: df_test.head()
```

```
Out[26]:
```

	Image	ID	Species	Breed
0	Abyssinian_201	1	1	1
1	Abyssinian_202	1	1	1
2	Abyssinian_204	1	1	1
3	Abyssinian_205	1	1	1
4	Abyssinian_206	1	1	1

```
In [27]: df_test['Species']=df_test['Species']-1
```

```
In [28]: print('No. of Cats:',df_test['Species'].value_counts()[0])
print('No. of Dogs:',df_test['Species'].value_counts()[1])
```

No. of Cats: 1182

No. of Dogs: 2486

```
In [29]: print('No. of Cat Species:',len(df_test.groupby('Species')['Breed'].unique()[0]))
print('No. of Dog Species:',len(df_test.groupby('Species')['Breed'].unique()[1]))
```

No. of Cat Species: 12

No. of Dog Species: 25

```
In [30]: test_tuple_df=np.array(df_test[['Image','Species']])
test_tuple_df

array([[ 'Abyssinian_201', 0],
```

```
Out[30]: [ 'Abyssinian_202', 0],
          [ 'Abyssinian_204', 0],
          ...,
          [ 'yorkshire_terrier_98', 1],
          [ 'yorkshire_terrier_99', 1],
          [ 'yorkshire_terrier_9', 1]], dtype=object)
```

## 'process\_test\_data(args...)' to process and prepare test data

```
In [31]: def process_test_data(df):

    data=list()

    for ele in df:

        image_name=ele[0]
        image_class=ele[1]
        image_path=os.path.join(all_images_path,image_name+'.jpg')
        img,_=process_image(image_path,None)
        data.append((image_name,img,None,image_class))

    print('No. of Elements processed:',len(data))

    return np.array(data)
```

```
In [32]: test_data=process_test_data(test_tuple_df)
```

No. of Elements processed: 3668

C:\Users\Asus\AppData\Local\Temp\ipykernel\_37160/2614160404.py:15: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray.

```
return np.array(data)
```

`test_data[0]-->image_name`

`test_data[1]-->img`

`test_data[2]-->None`

`test_data[3]-->image_class`

In [33]:

```
# Shuffling Training Data
np.random.shuffle(train_data)
```

## Splitting Data into Training and Validation Sets

In [34]:

```
val_data=train_data[:int(len(train_data)*0.2),:]
train_data=train_data[int(len(train_data)*0.2):,:]
```

In [35]:

```
x_train = []
y_class_train = []
y_box_train = []
x_validation = []
y_class_validation = []
y_box_validation = []
```

In [36]:

```
for img in train_data[:,1]:
    x_train.append(preprocess_input(img))

for box in train_data[:,2]:
    y_box_train.append(box)

for species in train_data[:,3]:
    y_class_train.append(species)
```

In [37]:

```
for img in val_data[:,1]:
    x_validation.append(preprocess_input(img))

for box in val_data[:,2]:
    y_box_validation.append(box)

for species in val_data[:,3]:
    y_class_validation.append(species)
```

## Preparing Test Set

In [38]:

```
x_test=[]
for img in test_data[:,1]:
    x_test.append(img)
```

## Casting Data to numpy array

In [39]:

```
x_train = np.array(x_train)
y_class_train = np.array(y_class_train)
y_box_train = np.array(y_box_train)
x_validation = np.array(x_validation)
y_class_validation = np.array(y_class_validation)
y_box_validation = np.array(y_box_validation)
x_test = np.array(x_test)
```

## Transfer Learning using Xception Model on ImageNet weights

In [40]:

```
base_model=Xception(weights='imagenet',include_top=False,input_shape=(224,224,3))
```

In [41]:

```
# Freezing 125 Layers out of total 132 Layers
print('Total no. of layers:',len(base_model.layers))
for x in range(125):
    base_model.layers[x].trainable=False
print('No. of frozen layers:',125)
```

Total no. of layers: 132

No. of frozen layers: 125

In [42]:

```
frozen_Xception=Model(inputs=[base_model.input],outputs=[base_model.output])
```

In [43]:

```
# Classification Output
classification_output = GlobalAveragePooling2D()(frozen_Xception.output)
classification_output = Dense(units=1, activation='sigmoid')(classification_output)

# Bounding Box Params Output
```

```
localization_output = Flatten()(frozen_Xception.output)
localization_output = Dense(units=4, activation='relu')(localization_output)

model = Model(inputs=[frozen_Xception.input], outputs=[classification_output, localization_output])

model.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #	Connected to
input_1 (InputLayer)	[None, 224, 224, 3 0 ]		[]
block1_conv1 (Conv2D)	(None, 111, 111, 32 864 )		['input_1[0][0]']
block1_conv1_bn (BatchNormalization)	(None, 111, 111, 32 128 )		['block1_conv1[0][0]']
block1_conv1_act (Activation)	(None, 111, 111, 32 0 )		['block1_conv1_bn[0][0]']
block1_conv2 (Conv2D)	(None, 109, 109, 64 18432 )		['block1_conv1_act[0][0]']
block1_conv2_bn (BatchNormalization)	(None, 109, 109, 64 256 )		['block1_conv2[0][0]']
block1_conv2_act (Activation)	(None, 109, 109, 64 0 )		['block1_conv2_bn[0][0]']
block2_sepconv1 (SeparableConv2D)	(None, 109, 109, 12 8768 )		['block1_conv2_act[0][0]']
block2_sepconv1_bn (BatchNormalization)	(None, 109, 109, 12 512 )		['block2_sepconv1[0][0]']
block2_sepconv2_act (Activation)	(None, 109, 109, 12 0 )		['block2_sepconv1_bn[0][0]']
block2_sepconv2 (SeparableConv2D)	(None, 109, 109, 12 17536 )		['block2_sepconv2_act[0][0]']

block2_sepconv2_bn (BatchNormalization)	(None, 109, 109, 128)	512	['block2_sepconv2[0][0]']
conv2d (Conv2D)	(None, 55, 55, 128)	8192	['block1_conv2_act[0][0]']
block2_pool (MaxPooling2D)	(None, 55, 55, 128)	0	['block2_sepconv2_bn[0][0]']
batch_normalization (BatchNormalization)	(None, 55, 55, 128)	512	['conv2d[0][0]']
add (Add)	(None, 55, 55, 128)	0	['block2_pool[0][0]', 'batch_normalization[0][0]']
block3_sepconv1_act (Activation)	(None, 55, 55, 128)	0	['add[0][0]']
block3_sepconv1 (SeparableConv2D)	(None, 55, 55, 256)	33920	['block3_sepconv1_act[0][0]']
block3_sepconv1_bn (BatchNormalization)	(None, 55, 55, 256)	1024	['block3_sepconv1[0][0]']
block3_sepconv2_act (Activation)	(None, 55, 55, 256)	0	['block3_sepconv1_bn[0][0]']
block3_sepconv2 (SeparableConv2D)	(None, 55, 55, 256)	67840	['block3_sepconv2_act[0][0]']
block3_sepconv2_bn (BatchNormalization)	(None, 55, 55, 256)	1024	['block3_sepconv2[0][0]']
conv2d_1 (Conv2D)	(None, 28, 28, 256)	32768	['add[0][0]']
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0	['block3_sepconv2_bn[0][0]']
batch_normalization_1 (BatchNormalization)	(None, 28, 28, 256)	1024	['conv2d_1[0][0]']
add_1 (Add)	(None, 28, 28, 256)	0	['block3_pool[0][0]', 'batch_normalization_1[0][0]']
block4_sepconv1_act (Activation)	(None, 28, 28, 256)	0	['add_1[0][0]']

block4_sepconv1 (SeparableConv 2D)	(None, 28, 28, 728)	188672	['block4_sepconv1_act[0][0]']
block4_sepconv1_bn (BatchNorma lization)	(None, 28, 28, 728)	2912	['block4_sepconv1[0][0]']
block4_sepconv2_act (Activatio n)	(None, 28, 28, 728)	0	['block4_sepconv1_bn[0][0]']
block4_sepconv2 (SeparableConv 2D)	(None, 28, 28, 728)	536536	['block4_sepconv2_act[0][0]']
block4_sepconv2_bn (BatchNorma lization)	(None, 28, 28, 728)	2912	['block4_sepconv2[0][0]']
conv2d_2 (Conv2D)	(None, 14, 14, 728)	186368	['add_1[0][0]']
block4_pool (MaxPooling2D)	(None, 14, 14, 728)	0	['block4_sepconv2_bn[0][0]']
batch_normalization_2 (BatchNo rmalization)	(None, 14, 14, 728)	2912	['conv2d_2[0][0]']
add_2 (Add)	(None, 14, 14, 728)	0	['block4_pool[0][0]', 'batch_normalization_2[0][0]']
block5_sepconv1_act (Activatio n)	(None, 14, 14, 728)	0	['add_2[0][0]']
block5_sepconv1 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block5_sepconv1_act[0][0]']
block5_sepconv1_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block5_sepconv1[0][0]']
block5_sepconv2_act (Activatio n)	(None, 14, 14, 728)	0	['block5_sepconv1_bn[0][0]']
block5_sepconv2 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block5_sepconv2_act[0][0]']
block5_sepconv2_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block5_sepconv2[0][0]']

block5_sepconv3_act (Activation)	(None, 14, 14, 728)	0	[ 'block5_sepconv2_bn[0][0]' ]
block5_sepconv3 (SeparableConv 2D)	(None, 14, 14, 728)	536536	[ 'block5_sepconv3_act[0][0]' ]
block5_sepconv3_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block5_sepconv3[0][0]' ]
add_3 (Add)	(None, 14, 14, 728)	0	[ 'block5_sepconv3_bn[0][0]', 'add_2[0][0]' ]
block6_sepconv1_act (Activation)	(None, 14, 14, 728)	0	[ 'add_3[0][0]' ]
block6_sepconv1 (SeparableConv 2D)	(None, 14, 14, 728)	536536	[ 'block6_sepconv1_act[0][0]' ]
block6_sepconv1_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block6_sepconv1[0][0]' ]
block6_sepconv2_act (Activation)	(None, 14, 14, 728)	0	[ 'block6_sepconv1_bn[0][0]' ]
block6_sepconv2 (SeparableConv 2D)	(None, 14, 14, 728)	536536	[ 'block6_sepconv2_act[0][0]' ]
block6_sepconv2_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block6_sepconv2[0][0]' ]
block6_sepconv3_act (Activation)	(None, 14, 14, 728)	0	[ 'block6_sepconv2_bn[0][0]' ]
block6_sepconv3 (SeparableConv 2D)	(None, 14, 14, 728)	536536	[ 'block6_sepconv3_act[0][0]' ]
block6_sepconv3_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block6_sepconv3[0][0]' ]
add_4 (Add)	(None, 14, 14, 728)	0	[ 'block6_sepconv3_bn[0][0]', 'add_3[0][0]' ]
block7_sepconv1_act (Activation)	(None, 14, 14, 728)	0	[ 'add_4[0][0]' ]

block7_sepconv1 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block7_sepconv1_act[0][0]']
block7_sepconv1_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block7_sepconv1[0][0]']
block7_sepconv2_act (Activatio n)	(None, 14, 14, 728)	0	['block7_sepconv1_bn[0][0]']
block7_sepconv2 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block7_sepconv2_act[0][0]']
block7_sepconv2_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block7_sepconv2[0][0]']
block7_sepconv3_act (Activatio n)	(None, 14, 14, 728)	0	['block7_sepconv2_bn[0][0]']
block7_sepconv3 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block7_sepconv3_act[0][0]']
block7_sepconv3_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block7_sepconv3[0][0]']
add_5 (Add)	(None, 14, 14, 728)	0	['block7_sepconv3_bn[0][0]', 'add_4[0][0]']
block8_sepconv1_act (Activatio n)	(None, 14, 14, 728)	0	['add_5[0][0]']
block8_sepconv1 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block8_sepconv1_act[0][0]']
block8_sepconv1_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block8_sepconv1[0][0]']
block8_sepconv2_act (Activatio n)	(None, 14, 14, 728)	0	['block8_sepconv1_bn[0][0]']
block8_sepconv2 (SeparableConv 2D)	(None, 14, 14, 728)	536536	['block8_sepconv2_act[0][0]']
block8_sepconv2_bn (BatchNorma lization)	(None, 14, 14, 728)	2912	['block8_sepconv2[0][0]']

lization)		
block8_sepconv3_act (Activation) (None, 14, 14, 728) 0 n)		['block8_sepconv2_bn[0][0]']
block8_sepconv3 (SeparableConv (None, 14, 14, 728) 536536 2D)		['block8_sepconv3_act[0][0]']
block8_sepconv3_bn (BatchNorma (None, 14, 14, 728) 2912 lization)		['block8_sepconv3[0][0]']
add_6 (Add) (None, 14, 14, 728) 0		['block8_sepconv3_bn[0][0]', 'add_5[0][0]']
block9_sepconv1_act (Activatio (None, 14, 14, 728) 0 n)		['add_6[0][0]']
block9_sepconv1 (SeparableConv (None, 14, 14, 728) 536536 2D)		['block9_sepconv1_act[0][0]']
block9_sepconv1_bn (BatchNorma (None, 14, 14, 728) 2912 lization)		['block9_sepconv1[0][0]']
block9_sepconv2_act (Activatio (None, 14, 14, 728) 0 n)		['block9_sepconv1_bn[0][0]']
block9_sepconv2 (SeparableConv (None, 14, 14, 728) 536536 2D)		['block9_sepconv2_act[0][0]']
block9_sepconv2_bn (BatchNorma (None, 14, 14, 728) 2912 lization)		['block9_sepconv2[0][0]']
block9_sepconv3_act (Activatio (None, 14, 14, 728) 0 n)		['block9_sepconv2_bn[0][0]']
block9_sepconv3 (SeparableConv (None, 14, 14, 728) 536536 2D)		['block9_sepconv3_act[0][0]']
block9_sepconv3_bn (BatchNorma (None, 14, 14, 728) 2912 lization)		['block9_sepconv3[0][0]']
add_7 (Add) (None, 14, 14, 728) 0		['block9_sepconv3_bn[0][0]', 'add_6[0][0]']

block10_sepconv1_act (Activation)	(None, 14, 14, 728)	0	[ 'add_7[0][0]' ]
block10_sepconv1 (SeparableConv2D)	(None, 14, 14, 728)	536536	[ 'block10_sepconv1_act[0][0]' ]
block10_sepconv1_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block10_sepconv1[0][0]' ]
block10_sepconv2_act (Activation)	(None, 14, 14, 728)	0	[ 'block10_sepconv1_bn[0][0]' ]
block10_sepconv2 (SeparableConv2D)	(None, 14, 14, 728)	536536	[ 'block10_sepconv2_act[0][0]' ]
block10_sepconv2_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block10_sepconv2[0][0]' ]
block10_sepconv3_act (Activation)	(None, 14, 14, 728)	0	[ 'block10_sepconv2_bn[0][0]' ]
block10_sepconv3 (SeparableConv2D)	(None, 14, 14, 728)	536536	[ 'block10_sepconv3_act[0][0]' ]
block10_sepconv3_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block10_sepconv3[0][0]' ]
add_8 (Add)	(None, 14, 14, 728)	0	[ 'block10_sepconv3_bn[0][0]', 'add_7[0][0]' ]
block11_sepconv1_act (Activation)	(None, 14, 14, 728)	0	[ 'add_8[0][0]' ]
block11_sepconv1 (SeparableConv2D)	(None, 14, 14, 728)	536536	[ 'block11_sepconv1_act[0][0]' ]
block11_sepconv1_bn (BatchNormalization)	(None, 14, 14, 728)	2912	[ 'block11_sepconv1[0][0]' ]
block11_sepconv2_act (Activation)	(None, 14, 14, 728)	0	[ 'block11_sepconv1_bn[0][0]' ]
block11_sepconv2 (SeparableConv2D)	(None, 14, 14, 728)	536536	[ 'block11_sepconv2_act[0][0]' ]

block11_sepconv2_bn (BatchNorm (None, 14, 14, 728) 2912 alization)		['block11_sepconv2[0][0]']
block11_sepconv3_act (Activati (None, 14, 14, 728) 0 on)		['block11_sepconv2_bn[0][0]']
block11_sepconv3 (SeparableCon (None, 14, 14, 728) 536536 v2D)		['block11_sepconv3_act[0][0]']
block11_sepconv3_bn (BatchNorm (None, 14, 14, 728) 2912 alization)		['block11_sepconv3[0][0]']
add_9 (Add) (None, 14, 14, 728) 0		['block11_sepconv3_bn[0][0]', 'add_8[0][0]']
block12_sepconv1_act (Activati (None, 14, 14, 728) 0 on)		['add_9[0][0]']
block12_sepconv1 (SeparableCon (None, 14, 14, 728) 536536 v2D)		['block12_sepconv1_act[0][0]']
block12_sepconv1_bn (BatchNorm (None, 14, 14, 728) 2912 alization)		['block12_sepconv1[0][0]']
block12_sepconv2_act (Activati (None, 14, 14, 728) 0 on)		['block12_sepconv1_bn[0][0]']
block12_sepconv2 (SeparableCon (None, 14, 14, 728) 536536 v2D)		['block12_sepconv2_act[0][0]']
block12_sepconv2_bn (BatchNorm (None, 14, 14, 728) 2912 alization)		['block12_sepconv2[0][0]']
block12_sepconv3_act (Activati (None, 14, 14, 728) 0 on)		['block12_sepconv2_bn[0][0]']
block12_sepconv3 (SeparableCon (None, 14, 14, 728) 536536 v2D)		['block12_sepconv3_act[0][0]']
block12_sepconv3_bn (BatchNorm (None, 14, 14, 728) 2912 alization)		['block12_sepconv3[0][0]']
add_10 (Add) (None, 14, 14, 728) 0		['block12_sepconv3_bn[0][0]', 'add_9[0][0]']

			'add_9[0][0]']
block13_sepconv1_act (Activation)	(None, 14, 14, 728) 0		['add_10[0][0]']
block13_sepconv1 (SeparableConv2D)	(None, 14, 14, 728) 536536		['block13_sepconv1_act[0][0]']
block13_sepconv1_bn (BatchNormalization)	(None, 14, 14, 728) 2912		['block13_sepconv1[0][0]']
block13_sepconv2_act (Activation)	(None, 14, 14, 728) 0		['block13_sepconv1_bn[0][0]']
block13_sepconv2 (SeparableConv2D)	(None, 14, 14, 1024) 752024		['block13_sepconv2_act[0][0]']
block13_sepconv2_bn (BatchNormalization)	(None, 14, 14, 1024) 4096		['block13_sepconv2[0][0]']
conv2d_3 (Conv2D)	(None, 7, 7, 1024) 745472		['add_10[0][0]']
block13_pool (MaxPooling2D)	(None, 7, 7, 1024) 0		['block13_sepconv2_bn[0][0]']
batch_normalization_3 (BatchNormalization)	(None, 7, 7, 1024) 4096		['conv2d_3[0][0]']
add_11 (Add)	(None, 7, 7, 1024) 0		['block13_pool[0][0]', 'batch_normalization_3[0][0]']
block14_sepconv1 (SeparableConv2D)	(None, 7, 7, 1536) 1582080		['add_11[0][0]']
block14_sepconv1_bn (BatchNormalization)	(None, 7, 7, 1536) 6144		['block14_sepconv1[0][0]']
block14_sepconv1_act (Activation)	(None, 7, 7, 1536) 0		['block14_sepconv1_bn[0][0]']
block14_sepconv2 (SeparableConv2D)	(None, 7, 7, 2048) 3159552		['block14_sepconv1_act[0][0]']
block14_sepconv2_bn (BatchNormalization)	(None, 7, 7, 2048) 8192		['block14_sepconv2[0][0]']

```

block14_sepconv2_act (Activation) (None, 7, 7, 2048) 0      ['block14_sepconv2_bn[0][0]']
on)

global_average_pooling2d (GlobalAveragePooling2D) (None, 2048) 0      ['block14_sepconv2_act[0][0]']

flatten (Flatten) (None, 100352) 0      ['block14_sepconv2_act[0][0]']

dense (Dense) (None, 1) 2049      ['global_average_pooling2d[0][0]']

dense_1 (Dense) (None, 4) 401412      ['flatten[0][0]']

=====
Total params: 21,264,941
Trainable params: 5,152,261
Non-trainable params: 16,112,680

```

---

## Fitting the Model

In [44]:

```

model.compile(optimizer='adam', metrics=['accuracy'], loss=['binary_crossentropy', 'mse'])
history=model.fit(x_train, [y_class_train, y_box_train], validation_data=(x_validation, [y_class_validation, y_box_validation])),ep
history

Epoch 1/20
92/92 [=====] - 433s 5s/step - loss: 991.3856 - dense_loss: 0.2650 - dense_1_loss: 991.1207 - dense_accuracy: 0.9231 - dense_1_accuracy: 0.8032 - val_loss: 534.2087 - val_dense_loss: 0.1049 - val_dense_1_loss: 534.1038 - val_dense_accuracy: 0.9837 - val_dense_1_accuracy: 0.8760
Epoch 2/20
92/92 [=====] - 413s 4s/step - loss: 154.1748 - dense_loss: 0.0995 - dense_1_loss: 154.0753 - dense_accuracy: 0.9813 - dense_1_accuracy: 0.9105 - val_loss: 266.8718 - val_dense_loss: 0.0687 - val_dense_1_loss: 266.8030 - val_dense_accuracy: 0.9932 - val_dense_1_accuracy: 0.8951
Epoch 3/20
92/92 [=====] - 402s 4s/step - loss: 94.8177 - dense_loss: 0.0827 - dense_1_loss: 94.7350 - dense_accuracy: 0.9823 - dense_1_accuracy: 0.9213 - val_loss: 143.6665 - val_dense_loss: 0.0593 - val_dense_1_loss: 143.6072 - val_dense_accuracy: 0.9905 - val_dense_1_accuracy: 0.9210
Epoch 4/20
92/92 [=====] - 408s 4s/step - loss: 71.0263 - dense_loss: 0.0681 - dense_1_loss: 70.9582 - dense_accuracy: 0.9857 - dense_1_accuracy: 0.9282 - val_loss: 123.4953 - val_dense_loss: 0.0523 - val_dense_1_loss: 123.4429 - val_dense_accuracy: 0.9905 - val_dense_1_accuracy: 0.9196

```

Epoch 5/20  
92/92 [=====] - 410s 4s/step - loss: 70.5779 - dense\_loss: 0.0622 - dense\_1\_loss: 70.5157 - dense\_accuracy: 0.9823 - dense\_1\_accuracy: 0.9367 - val\_loss: 174.5141 - val\_dense\_loss: 0.0489 - val\_dense\_1\_loss: 174.4653 - val\_dense\_accuracy: 0.9891 - val\_dense\_1\_accuracy: 0.9251  
Epoch 6/20  
92/92 [=====] - 383s 4s/step - loss: 48.7992 - dense\_loss: 0.0612 - dense\_1\_loss: 48.7380 - dense\_accuracy: 0.9850 - dense\_1\_accuracy: 0.9404 - val\_loss: 109.3161 - val\_dense\_loss: 0.0458 - val\_dense\_1\_loss: 109.2703 - val\_dense\_accuracy: 0.9918 - val\_dense\_1\_accuracy: 0.9251  
Epoch 7/20  
92/92 [=====] - 412s 4s/step - loss: 47.0993 - dense\_loss: 0.0508 - dense\_1\_loss: 47.0485 - dense\_accuracy: 0.9901 - dense\_1\_accuracy: 0.9435 - val\_loss: 167.4933 - val\_dense\_loss: 0.0374 - val\_dense\_1\_loss: 167.4559 - val\_dense\_accuracy: 0.9918 - val\_dense\_1\_accuracy: 0.9210  
Epoch 8/20  
92/92 [=====] - 393s 4s/step - loss: 36.3960 - dense\_loss: 0.0527 - dense\_1\_loss: 36.3432 - dense\_accuracy: 0.9840 - dense\_1\_accuracy: 0.9544 - val\_loss: 130.7314 - val\_dense\_loss: 0.0401 - val\_dense\_1\_loss: 130.6913 - val\_dense\_accuracy: 0.9905 - val\_dense\_1\_accuracy: 0.9087  
Epoch 9/20  
92/92 [=====] - 402s 4s/step - loss: 42.6278 - dense\_loss: 0.0477 - dense\_1\_loss: 42.5801 - dense\_accuracy: 0.9898 - dense\_1\_accuracy: 0.9537 - val\_loss: 120.6181 - val\_dense\_loss: 0.0356 - val\_dense\_1\_loss: 120.5824 - val\_dense\_accuracy: 0.9905 - val\_dense\_1\_accuracy: 0.9210  
Epoch 10/20  
92/92 [=====] - 435s 5s/step - loss: 32.9885 - dense\_loss: 0.0445 - dense\_1\_loss: 32.9440 - dense\_accuracy: 0.9888 - dense\_1\_accuracy: 0.9687 - val\_loss: 130.0070 - val\_dense\_loss: 0.0424 - val\_dense\_1\_loss: 129.9647 - val\_dense\_accuracy: 0.9918 - val\_dense\_1\_accuracy: 0.9169  
Epoch 11/20  
92/92 [=====] - 385s 4s/step - loss: 27.2425 - dense\_loss: 0.0402 - dense\_1\_loss: 27.2022 - dense\_accuracy: 0.9915 - dense\_1\_accuracy: 0.9663 - val\_loss: 133.1241 - val\_dense\_loss: 0.0371 - val\_dense\_1\_loss: 133.0870 - val\_dense\_accuracy: 0.9877 - val\_dense\_1\_accuracy: 0.9128  
Epoch 12/20  
92/92 [=====] - 403s 4s/step - loss: 27.1508 - dense\_loss: 0.0426 - dense\_1\_loss: 27.1082 - dense\_accuracy: 0.9908 - dense\_1\_accuracy: 0.9714 - val\_loss: 116.0304 - val\_dense\_loss: 0.0352 - val\_dense\_1\_loss: 115.9952 - val\_dense\_accuracy: 0.9918 - val\_dense\_1\_accuracy: 0.9155  
Epoch 13/20  
92/92 [=====] - 399s 4s/step - loss: 25.4699 - dense\_loss: 0.0387 - dense\_1\_loss: 25.4312 - dense\_accuracy: 0.9905 - dense\_1\_accuracy: 0.9694 - val\_loss: 94.7238 - val\_dense\_loss: 0.0365 - val\_dense\_1\_loss: 94.6873 - val\_dense\_accuracy: 0.9905 - val\_dense\_1\_accuracy: 0.9251  
Epoch 14/20  
92/92 [=====] - 364s 4s/step - loss: 24.7988 - dense\_loss: 0.0358 - dense\_1\_loss: 24.7630 - dense\_accuracy: 0.9911 - dense\_1\_accuracy: 0.9741 - val\_loss: 90.3766 - val\_dense\_loss: 0.0353 - val\_dense\_1\_loss: 90.3413 - val\_dense\_accuracy: 0.9891 - val\_dense\_1\_accuracy: 0.9223  
Epoch 15/20  
92/92 [=====] - 420s 5s/step - loss: 22.4888 - dense\_loss: 0.0375 - dense\_1\_loss: 22.4513 - dense\_accuracy: 0.9901 - dense\_1\_accuracy: 0.9741 - val\_loss: 101.4216 - val\_dense\_loss: 0.0359 - val\_dense\_1\_loss: 101.3857 - val\_dense\_accuracy: 0.9891 - val\_dense\_1\_accuracy: 0.9210

```

Epoch 16/20
92/92 [=====] - 370s 4s/step - loss: 22.3997 - dense_loss: 0.0386 - dense_1_loss: 22.3611 - dense_accuracy: 0.9901 - dense_1_accuracy: 0.9704 - val_loss: 105.7388 - val_dense_loss: 0.0334 - val_dense_1_loss: 105.7055 - val_dense_accuracy: 0.9918 - val_dense_1_accuracy: 0.9237
Epoch 17/20
92/92 [=====] - 408s 4s/step - loss: 24.1302 - dense_loss: 0.0343 - dense_1_loss: 24.0959 - dense_accuracy: 0.9901 - dense_1_accuracy: 0.9741 - val_loss: 98.4875 - val_dense_loss: 0.0312 - val_dense_1_loss: 98.4564 - val_dense_accuracy: 0.9918 - val_dense_1_accuracy: 0.9251
Epoch 18/20
92/92 [=====] - 385s 4s/step - loss: 17.9329 - dense_loss: 0.0367 - dense_1_loss: 17.8962 - dense_accuracy: 0.9891 - dense_1_accuracy: 0.9762 - val_loss: 85.9742 - val_dense_loss: 0.0379 - val_dense_1_loss: 85.9363 - val_dense_accuracy: 0.9905 - val_dense_1_accuracy: 0.9210
Epoch 19/20
92/92 [=====] - 479s 5s/step - loss: 17.0046 - dense_loss: 0.0340 - dense_1_loss: 16.9706 - dense_accuracy: 0.9905 - dense_1_accuracy: 0.9765 - val_loss: 105.6739 - val_dense_loss: 0.0368 - val_dense_1_loss: 105.6371 - val_dense_accuracy: 0.9891 - val_dense_1_accuracy: 0.9169
Epoch 20/20
92/92 [=====] - 488s 5s/step - loss: 17.0071 - dense_loss: 0.0308 - dense_1_loss: 16.9762 - dense_accuracy: 0.9922 - dense_1_accuracy: 0.9700 - val_loss: 111.0330 - val_dense_loss: 0.0313 - val_dense_1_loss: 111.0016 - val_dense_accuracy: 0.9905 - val_dense_1_accuracy: 0.9169
Out[44]:
<keras.callbacks.History at 0x293db9b5700>

```

In [45]: `history.history.keys()`

Out[45]: `dict_keys(['loss', 'dense_loss', 'dense_1_loss', 'dense_accuracy', 'dense_1_accuracy', 'val_loss', 'val_dense_loss', 'val_dense_1_loss', 'val_dense_accuracy', 'val_dense_1_accuracy'])`

## Training History Plot for Classification

In [46]:

```

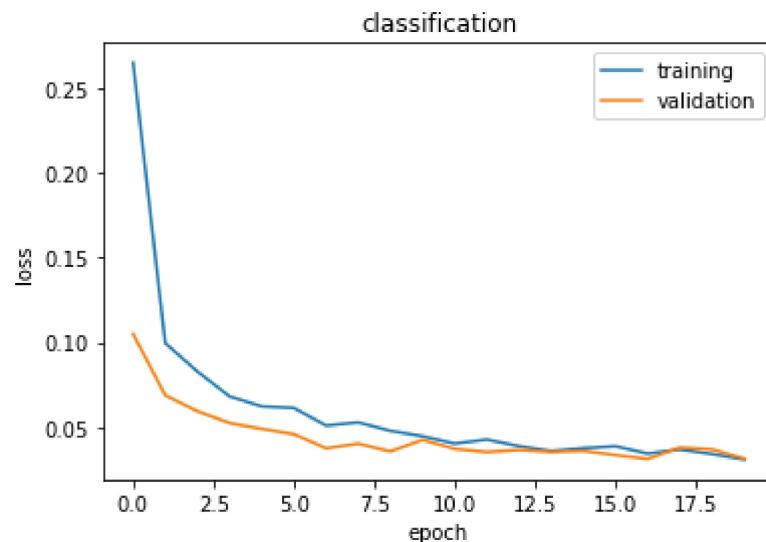
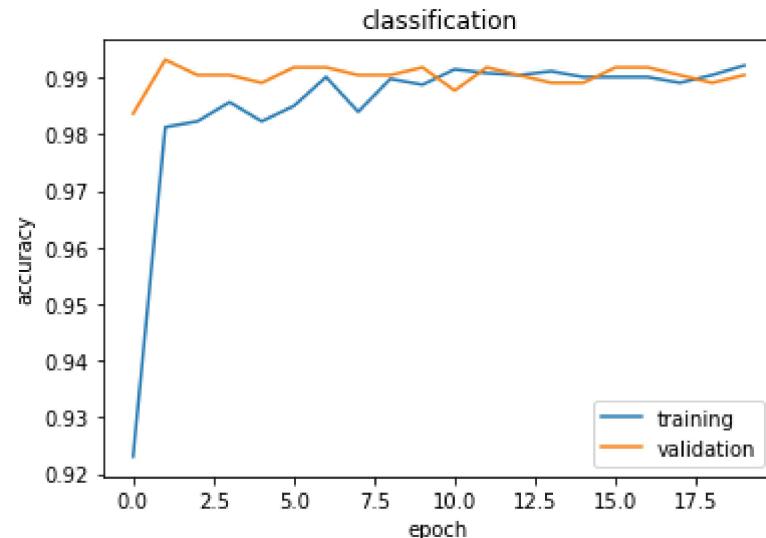
def plot_training_history(history, model):
    plt.plot(history.history['dense_accuracy'])
    plt.plot(history.history['val_dense_accuracy'])
    plt.title('classification')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['training', 'validation'], loc='best')
    plt.show()

    plt.plot(history.history['dense_loss'])
    plt.plot(history.history['val_dense_loss'])
    plt.title('classification')

```

```
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['training', 'validation'], loc='best')
plt.show()

plot_training_history(history, model)
```



Training and Validation Accuracy are very close to each other for classification task, indicating ideal convergence

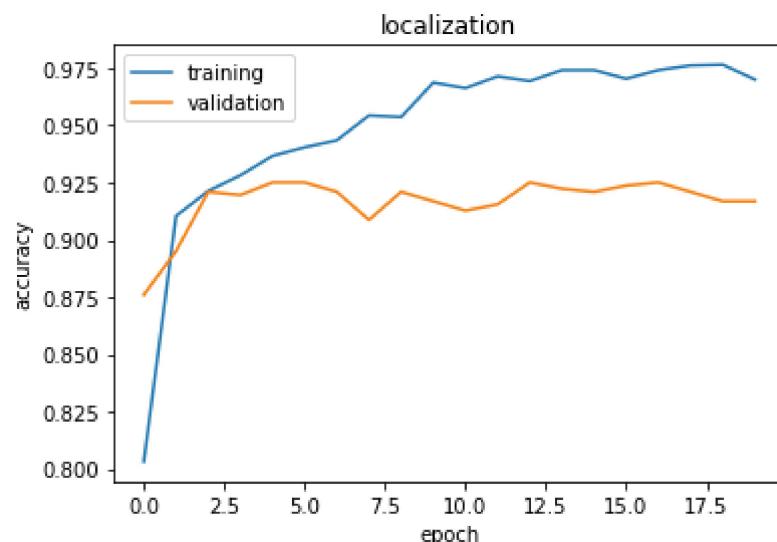
## Training History Plot for Localization

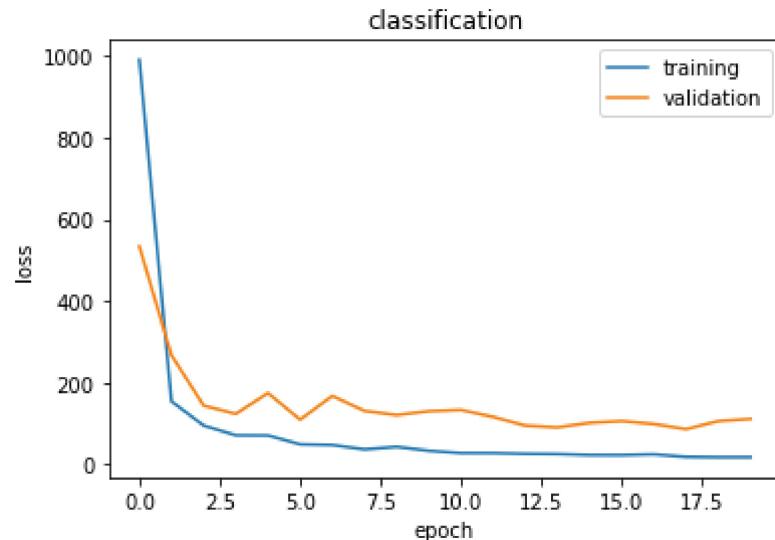
In [47]:

```
def plot_training_history(history, model):
    plt.plot(history.history['dense_1_accuracy'])
    plt.plot(history.history['val_dense_1_accuracy'])
    plt.title('localization')
    plt.ylabel('accuracy')
    plt.xlabel('epoch')
    plt.legend(['training', 'validation'], loc='best')
    plt.show()

    plt.plot(history.history['dense_1_loss'])
    plt.plot(history.history['val_dense_1_loss'])
    plt.title('classification')
    plt.ylabel('loss')
    plt.xlabel('epoch')
    plt.legend(['training', 'validation'], loc='best')
    plt.show()

plot_training_history(history, model)
```





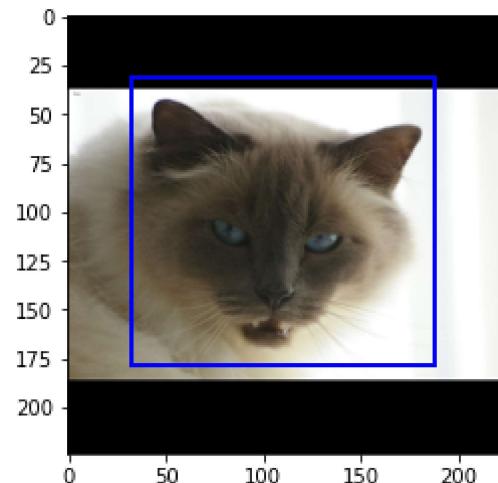
Training and Validation Accuracy are less close to each other for localization task, indicating some degree of overfitting

## Testing Model on Unseen Data

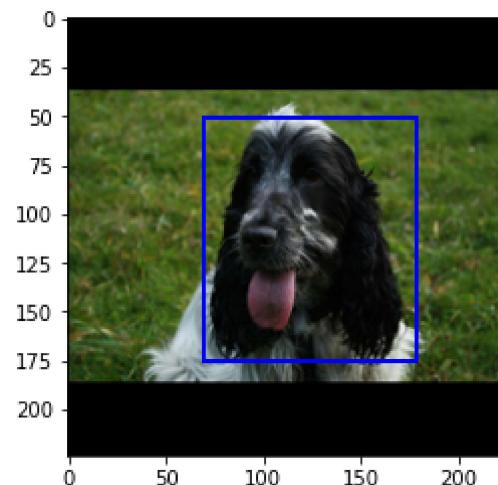
In [56]:

```
for i in range(6):
    n=np.random.randint(len(x_test))
    pred = model.predict(np.array([preprocess_input(x_test[n])]))
    if pred[0][0]>0.5:
        print('it is a dog')
    else:
        print('it is a cat')
    image_with_bndbox(x_test[n],pred_bndbox=Bounding_Box(*pred[1][0]))
```

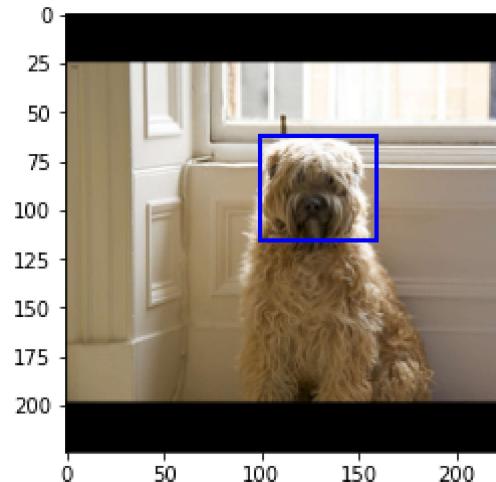
it is a cat



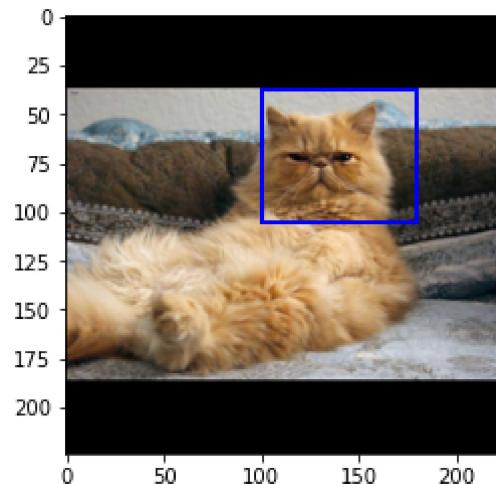
it is a dog



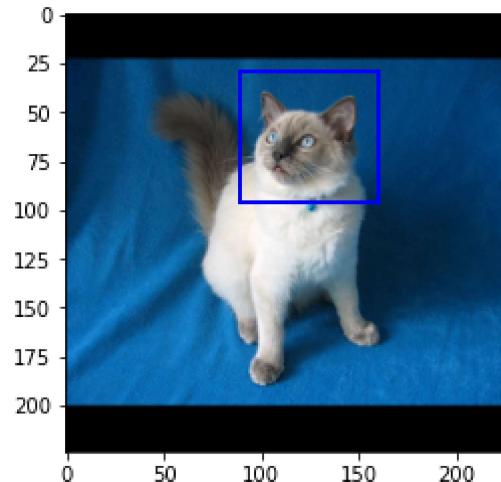
it is a dog



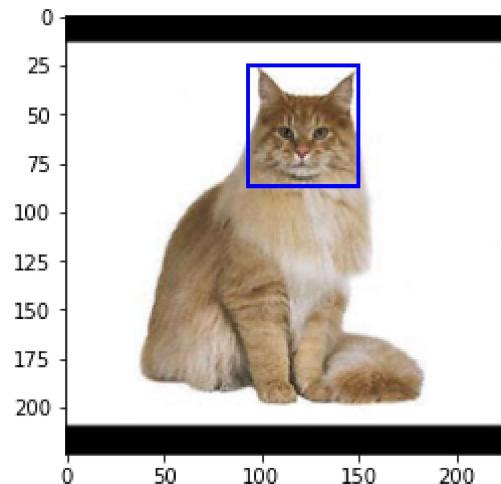
it is a cat



it is a cat



it is a cat



**It can be clearly observed that the model shows promising results even on unseen data**

**Further Suggestions to improve model performance:**

# More varied data can be added to tackle the problem of overfitting

# Data Augmentation can be performed

# Freezing of different no. of layers can be tried for fine tuning

# Different coordinate setup can be tried for localization problem

In [ ]: